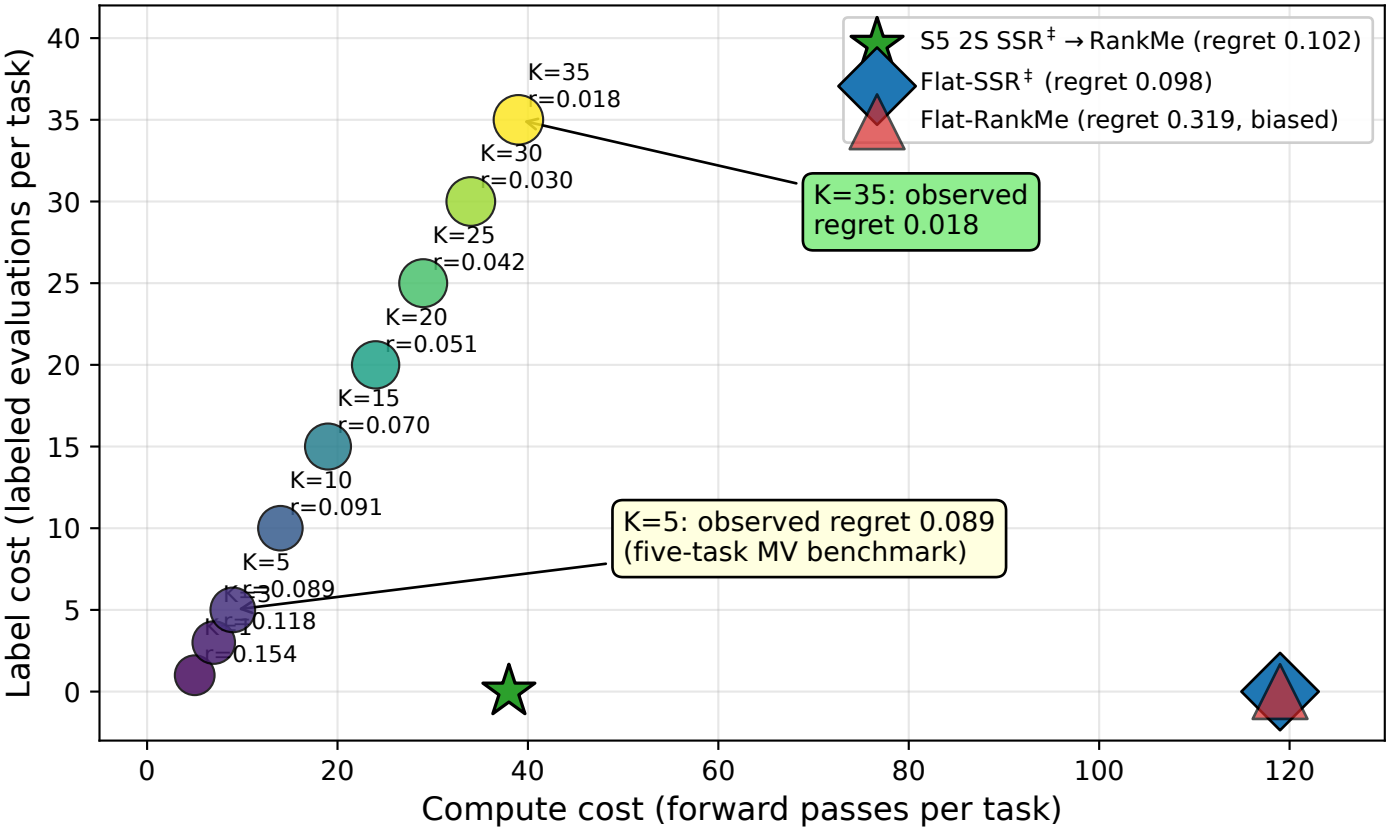
Cost--quality trade-off: compute vs label budget axes
Bubble size ∝ regret reduction; mixed pipeline (H-score → RankMe) sweep K=1..35
S5 2S SSR‡ → RankMe (regret 0.102)
Flat-SSR‡ (regret 0.098)
Flat-RankMe (regret 0.319, biased)
K=35: observed regret 0.018
K=5: observed regret 0.089 (five-task MV benchmark)
K=35 r=0.018
K=30 r=0.030
K=25 r=0.042
K=20 r=0.051
K=15 r=0.070
K=10 r=0.091
K=5
r=0.089
r=0.118
r=0.154
Label cost (labeled evaluations per task)
Compute cost (forward passes per task)
0
5
10
15
20
25
30
35
40
0
20
40
60
80
100
120